

Developer Operations 2024 - Assignment 2

Student Number: 20104119

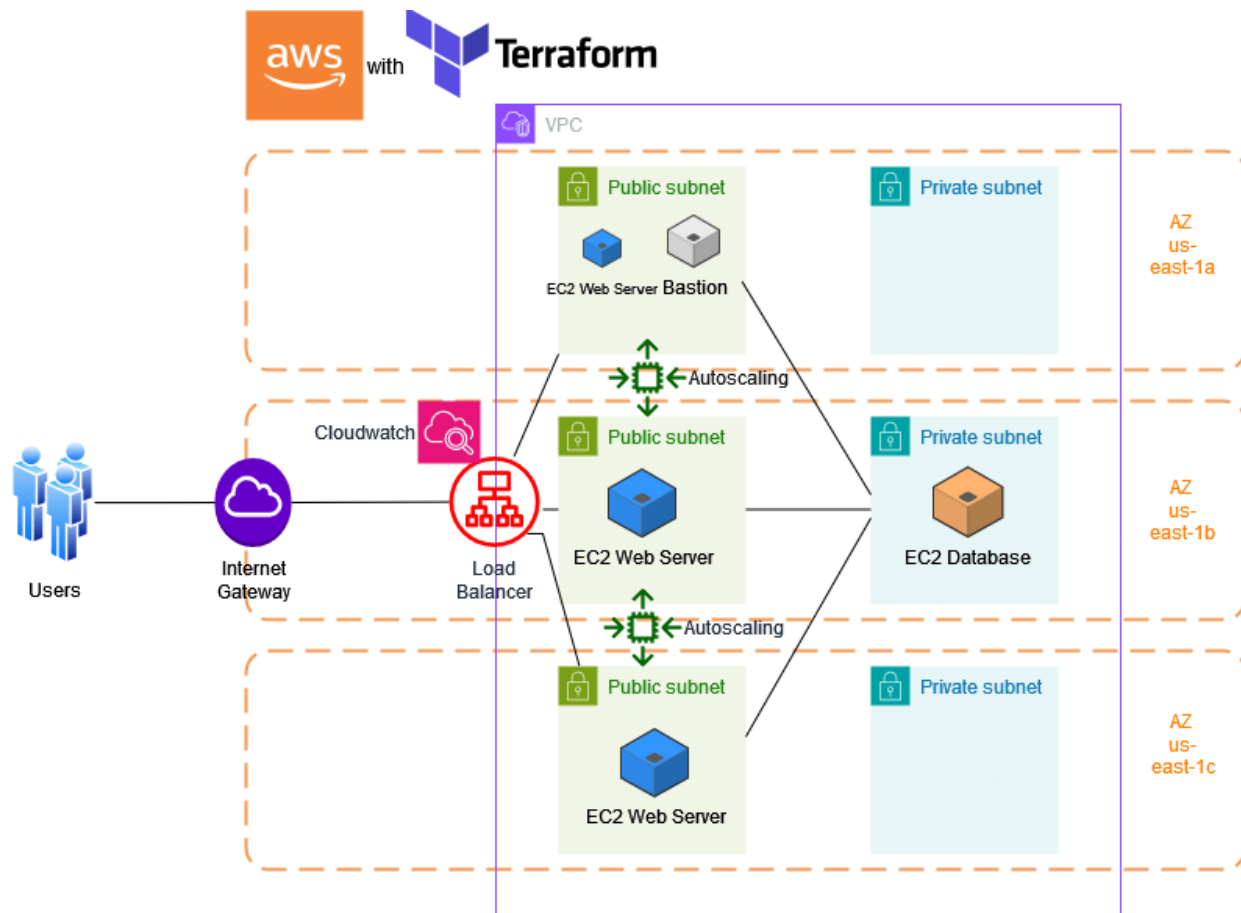
Student Name: Marty Rath

Project URL: <https://github.com/MartyRath/aws-web-app-autoscaling>

Table of Contents

Architecture overview	3
Step 1: Master instance	4
Screenshot 1: Master instance	4
Step 2: Custom AMI	6
Screenshot 2: Custom AMI	6
Step 3a: VPC and subnets	7
Screenshot 3a: Public/Private subnet AZ distribution	7
Screenshot 3a - 1: Public route table using IGW	8
Screenshot 3a - 2: Private route table using Nat gateway	8
Screenshot 3a -3: Nat gateway	8
Step 3b: Security groups (rules listed)	9
Screenshot 3b: Mongo instance private IP showing NAT gateway IP	9
Step 4a: Load balancer	11
Screenshot 4a: Load Balancer (round robin)	11
Screenshot 4a - 1: Target group	11
Screenshot 4a - 2: Listeners. HTTP and HTTPS	12
Step 4b: Launch template	13
Screenshot 4b: Launch Template	13
Step 4c: Auto-scaling group	14
Screenshot 4c: Auto-scaling Group (shows Launch Template, subnet AZs, load balancer)	14
Step 5: Scaling Policies, CloudWatch Alarms	16
Screenshot 5: Scaling policies	16
Screenshot 5 - 1: Auto-scaled instances	16
Step 6: Test Traffic	17
Screenshot 6: Test traffic generation script	17
Step 7: Load Distribution	18
Screenshot 7: Load Distribution (browser)	18
Step 8: Custom metrics	19
Screenshot 8: Custom metrics in CloudWatch	19
Additional Functionality	21
- Screenshot: HTTPS set up:	21

Architecture overview



The architecture is designed to ensure high availability, scalability and automated management of the web application running on the EC2 Web Servers.

Key points are:

- The web application. This was configured on a master instance, from which a custom Amazon Machine Image (AMI) was built. The custom AMI is the template for scaled instances.
- The Virtual Private Cloud (VPC) provides an isolated network environment for deploying the application. The web application runs on public subnets to enable internet accessibility via the Internet Gateway. The database instance is run on a private subnet for security.
- Traffic distribution across the web application instances is managed by the load balancer. This is to maintain high quality of service, dynamically scaling based on demand.
- Template instances of the web application from the custom AMI are dynamically launched or destroyed by the auto-scaling.
- CloudWatch alarms monitor instance metrics, informing the scaling policies to adjust instance count as needed, ensuring consistent performance.

The remainder of this report will go into detail on how this architecture was achieved.

Step 1: Master instance

See 01-master-instance.tf

The Terraform resource `aws_ami` was used to retrieve the most recent Amazon Linux x86_64 AMI for the master instance. The AMI was searched for via `aws_ami`, with results filtered using globbing:

```
filter {
  name = "name"
  values = ["al2023-ami-2023*x86_64"]
}
```

`Aws_instance` was used to create the master instance in a public subnet. Other configurations include, a web server security group (discussed Step 3b), a key pair for SSH access, and an IAM instance profile (`LabInstanceProfile`) in order to push custom metrics to CloudWatch.

User data included the script `deploy_app.sh`. In this script, it uploaded the `start_app.sh` script to install the Playtime web application onto the instance. The script also installed dependencies for the app and created a static web page with the EC2 ID at `id.html`.

In the end, the node app was installed, along with custom metrics pushed to CloudWatch via SSH when attempts to automate this through `user_data` scripts failed.

This was done to create a new custom AMI with the web app properly installed.

Screenshot 1: Master instance

i-0a0dfb20c6607b06f (Main Web Server)

▼ Instance summary Info

Instance ID

i-0a0dfb20c6607b06f (Main Web Server)

IPv6 address

—

Hostname type

IP name: ip-10-0-101-218.ec2.internal

Answer private resource DNS name

—

Auto-assigned IP address

18.232.102.136 [Public IP]

IAM Role

LabRole

Public IPv4 address

18.232.102.136 | open address

Instance state

Running

Private IP DNS name (IPv4 only)

ip-10-0-101-218.ec2.internal

Instance type

t2.nano

VPC ID

vpc-09296590b3eea2b74 (myVPC)

Private IPv4 addresses

10.0.101.218

Public IPv4 DNS

ec2-18-232-102-136.compute-1.amazonaws.com | open address

Elastic IP addresses

—

AWS Compute Optimizer finding

Opt-in to AWS Compute Optimizer for recommendations. |

Subnet ID

subnet-0487e3d73a1e57bf7 (public_subnet_1)

Step 2: Custom AMI

See 01-master-instance.tf

The custom AMI was based on the master instance, and so shared the configurations discussed above. This was achieved by using the Terraform resource: `aws_ami_from_instance`

The custom AMI is tagged as, “Custom Web Server AMI”.

The AMI is linked to the master instance via the instance id:

```
source_instance_id = aws_instance.main_web_server.id
```

Screenshot 2: Custom AMI

<

Step 3a: VPC and subnets

See 02-vpc-security-groups.tf

A Terraform module was used to create the VPC. The documentation can be found here:

<https://registry.terraform.io/modules/terraform-aws-modules/vpc/aws/latest>

Along with the VPC, it was specified to create three private and public subnets, in the availability zones: us-east-1a, us-east-1b, us-east-1c.

Screenshot 3a: Public/Private subnet AZ distribution

Name	Subn...	State	VPC	IPv4 CIDR	IPv6 ...	Avail...	Availability Zone
public_subnet_3	subnet-...	✔ Available	vpc-092...	10.0.103.0/24	–	249	us-east-1c
public_subnet_2	subnet-...	✔ Available	vpc-092...	10.0.102.0/24	–	250	us-east-1b
public_subnet_1	subnet-...	✔ Available	vpc-092...	10.0.101.0/24	–	248	us-east-1a
private_subnet_3	subnet-...	✔ Available	vpc-092...	10.0.3.0/24	–	251	us-east-1c
private_subnet_2	subnet-...	✔ Available	vpc-092...	10.0.2.0/24	–	251	us-east-1b
private_subnet_1	subnet-...	✔ Available	vpc-092...	10.0.1.0/24	–	250	us-east-1a

Other customisations to this module were to:

Auto-assign public IPv4 address for instances launched in public subnets

```
map_public_ip_on_launch = true
```

Enable a NAT gateway in just one availability zone for private subnets.

```
enable_nat_gateway = true
```

```
single_nat_gateway = true
```

And to disable the creation of a default security group.

```
manage_default_security_group = false
```

By default, this module sets up custom route tables, an internet gateway for the public subnets, and a network access control list.

The default network ACL allows traffic from any source, but the security groups (step 3b) take precedence over this rule.

Screenshot 3a - 1: Public route table using IGW

rtb-0a08a5226144eca3c / myVPC-public

Details

Routes

Subnet associations

Edge associations

Route propagation

Tags

Routes (2)			
<input type="text" value="Filter routes"/>			
Destination	Target	Status	Propagated
0.0.0.0/0	igw-08bad6073dee5e5d4	Active	No
10.0.0.0/16	local	Active	No

Screenshot 3a - 2: Private route table using Nat gateway

rtb-026e5dace3348c748 / myVPC-private

Details

Routes

Subnet associations

Edge associations

Route propagation

Tags

Routes (2)			
<input type="text" value="Filter routes"/>			
Destination	Target	Status	Propagated
0.0.0.0/0	nat-02a2142f1b33999a4	Active	No
10.0.0.0/16	local	Active	No

Screenshot 3a -3: Nat gateway

	Name	NAT gateway ID	Connectivity...	State	State message	Primary public I...	Primary private I...
	myVPC-us-east-1a	nat-02a2142f1b33999a4	Public	Available	-	54.146.223.80	10.0.101.243

nat-02a2142f1b33999a4 / myVPC-us-east-1a

Details

Secondary IPv4 addresses

Monitoring

Tags

Details			
NAT gateway ID nat-02a2142f1b33999a4	Connectivity type Public	State Available	State message -
NAT gateway ARN arn:aws:ec2:us-east-1:701357344768:natgateway/nat-02a2142f1b33999a4	Primary public IPv4 address 54.146.223.80	Primary private IPv4 address 10.0.101.243	Primary network interface ID eni-0878007c79c68848f
VPC vpc-09296590b3eea2b74 / myVPC	Subnet subnet-0487e3d73a1e57bf7 / public_subnet_1	Created Sunday, 28 April 2024 at 12:21:18 GMT+1	Deleted -

Step 3b: Security groups (rules listed)

See 02-vpc-security-groups.tf

I used a Terraform resource to create multiple security groups.

Web server security group:

Ingress rules to allow:

- SSH traffic from any IP
- HTTP traffic from any IP
- HTTPS traffic from any IP
- Custom TCP traffic from any IP

Egress rule:

- All outbound traffic from anywhere

Bastion Security Group

Ingress rules to allow:

- SSH traffic from any IP

Egress rule to allow:

- SSH traffic from any IP

Mongo Security Group

Ingress rules to allow:

- SSH traffic from the Bastion
- TCP traffic from the web server/node app.

Egress rule:

- All outbound traffic from anywhere







Screenshot 3b: Mongo instance private IP showing NAT gateway IP

i-0d21e0964073650e7 (Mongo Database)

The screenshot displays the AWS Management Console for an EC2 instance. The instance is named 'i-0d21e0964073650e7 (Mongo Database)' and is in the 'Running' state. The private IP address is 10.0.1.220. A terminal window at the bottom shows a command being executed: `wget http://ipinfo.io/ip -q0 - 52.22.28.19`.

Details	Status and alarms New	Monitoring	Security	Networking	Storage	Tags
▼ Instance summary Info						
Instance ID	i-0d21e0964073650e7 (Mongo Database)					
IPv6 address	-					
Hostname type	IP name: ip-10-0-1-220.ec2.internal					
Public IPv4 address	-					
Instance state	✓ Running					
Private IP DNS name (IPv4 only)	ip-10-0-1-220.ec2.internal					
Private IPv4 addresses	10.0.1.220					
Public IPv4 DNS	-					

```
[ec2-user@ip-10-0-1-220 ~]$ wget http://ipinfo.io/ip -q0 - 52.22.28.19[ec2-user@ip-10-0-1-220 ~]$
```

Details			
<div>NAT gateway ID</div> <div> nat-09cee071c9fb8ac5f</div>	<div>Connectivity type</div> <div>Public</div>	<div>State</div> <div> Available</div>	<div>State message</div> <div>-</div>
<div>NAT gateway ARN</div> <div> arn:aws:ec2:us-east-1:701357344768:natgateway/nat-09cee071c9fb8ac5f</div>	<div>Primary public IPv4 address</div> <div>52.22.28.19</div>	<div>Primary private IPv4 address</div> <div> 10.0.101.244</div>	<div>Primary network interface ID</div> <div>eni-09de0f221cb6d254b </div>
<div>VPC</div> <div>vpc-0733dc534755db63a / myVPC</div>	<div>Subnet</div> <div>subnet-051e3495709a86c45 / public_subnet_1</div>	<div>Created</div> <div> Sunday, 28 April 2024 at 19:01:46 GMT+1</div>	<div>Deleted</div> <div>-</div>

Step 4a: Load balancer

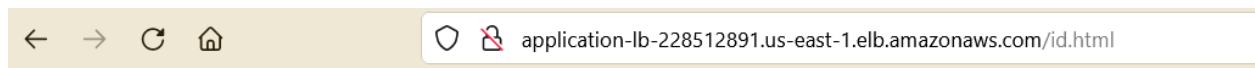
See 03-load-balancer.tf

An application Load Balancer (ALB) along with a target group and listeners for handling HTTP traffic using Terraform.


The target group was then updated to port 3000 for traffic from the Node app.

The ALB distributes incoming requests to the instances registered with the target group based on the configured listeners and routing rules.

Screenshot 4a: Load Balancer (round robin)



Instance ID: i-060e307915ea3b3b7



Instance ID: i-002df13b55c7c8fa1

Screenshot 4a - 1: Target group

Target group: web-server-tg

Details			
arn:aws:elasticloadbalancing:us-east-1:701357344768:targetgroup/web-server-tg/254abb9cbac2b592			
Target type	Protocol : Port	Protocol version	VPC
Instance	HTTP: 3000	HTTP1	vpc-0733dc534755db63a
IP address type	Load balancer		
IPv4	application-lb		

Screenshot 4a - 2: Listeners. HTTP and HTTPS

Load balancer: application-lb

Listeners and rules (2) Info							Manage rules	Manage listener	Add listener
<input type="text" value="Filter listeners"/>							< 1 >		
<input type="checkbox"/>	Protocol:Port	Default action	Rules	ARN	Security policy	Default SSL/TLS certificate			
<input type="checkbox"/>	HTTP:80	Forward to target group <ul style="list-style-type: none">web-server-tg : 1 (100%)Group-level stickiness: Off	1 rule	ARN	Not applicable	Not applicable			
<input type="checkbox"/>	HTTPS:443	Forward to target group <ul style="list-style-type: none">web-server-tg : 1 (100%)Group-level stickiness: Off	1 rule	ARN	ELBSecurityPolicy-2016-08	application-lb-1203723242.us...			

Step 4b: Launch template

See 04-autoscaling.tf

The launch template is created from the custom AMI created called localMongo. It has the web server security group and a key pair for SSH access.

As custom metrics will be run every minute, it also enables detailed monitoring, and for permissions to push custom metrics to CloudWatch has the IAM instance profile, LabInstanceProfile.

Then, its user data was used to start the node app with:

```
su - ec2-user -c 'cd playtime; npm run start'
```

Screenshot 4b: Launch Template

terraform-20240428135845979200000006 (lt-01dbb85811aad9381)

Version

1 (Default)

Description

-

Date created

2024-04-28T13:58:46.000Z

Created by

arn:aws:sts::701357344768:assumed-role/voclabs/user3015540=Marty_Rath

Instance details

Storage

Resource tags

Network interfaces

Advanced details

AMI ID

ami-05ca988622607f072

Instance type

t2.nano

Availability Zone

-

Key pair name

firstLabKey

Security groups

-

Security group IDs

sg-0a881fb62f41e098

Step 4c: Auto-scaling group

See 04-autoscaling.tf


The web server auto-scaling group (ASG) uses the launch template and is linked to the load balancer.




The ASG ensures that the number of instances dynamically scales based on CPU utilisation, with a maximum size of 3 instances, a minimum size of 1 instance, and a desired capacity of 1 instance.

The ASG is configured to use 1-minute granularity for metrics within CloudWatch, allowing for detailed monitoring and scaling actions.

Additionally, each autoscaled instance is tagged with a timestamp for identification.

Screenshot 4c: Auto-scaling Group (shows Launch Template, subnet AZs, load balancer)

Group details			Edit
Auto Scaling group name web-server-asg	Desired capacity 2	Desired capacity type Units (number of instances)	Amazon Resource Name (ARN)  arn:aws:autoscaling:us-east-1:701357344768:autoScalingGroup:e33219dd-5d93-4381-b060-5dfbad39047a:autoScalingGroupName/web-server-asg
Date created Sun Apr 28 2024 14:58:48 GMT+0100 (Irish Standard Time)	Minimum capacity 1	Status -	
	Maximum capacity 3		

Launch template			Edit
Launch template  lt-01dbb85811aad9381 terraform-20240428135845979200000006	AMI ID  ami-05ca988622607f072	Instance type t2.nano	Owner arn:aws:sts::701357344768:assumed-role/voclabs/user3015540=Marty_Rath Create time Sun Apr 28 2024 14:58:46 GMT+0100 (Irish Standard Time) Request Spot Instances No
Version Latest	Security groups -	Security group IDs  sg-0a881fbf62f41e098	
Description -	Storage (volumes) -	Key pair name firstLabKey	

Network

[Edit](#)

Availability Zones

us-east-1a, us-east-1b, us-east-1c

Subnet ID

subnet-093453e05f92b3ebe,
subnet-08507d23034124a12,
subnet-05f2509e216652e69

Load balancing

[Edit](#)

Load balancer target groups

[web-server-tg](#)

Classic Load Balancers

-

Health checks

[Edit](#)

Health check type

EC2

Health check grace period

30

Step 5: Scaling Policies, CloudWatch Alarms

See 04-autoscaling.tf

There are two autoscaling policies in place, designed to scale out on high CPU, and scale in on low CPU.

These policies are linked to the high_CPU_alarm and low_CPU_alarm CloudWatch alarms. These alarms are configured to notify the autoscaling policies of high/low CPU after two consequent checks of CPU utilisation , which are done every minute.

The high_CPU_alarm is triggered when CPU utilisation is on average above 40%, and the low_CPU_alarm is triggered when CPU utilisation is on average below 20%.

Screenshot 5: Scaling policies

Auto Scaling group: web-server-asg

scale-in-low-CPU-asp

Simple scaling

Enabled

low-CPU-alarm

breaches the alarm threshold: CPUUtilization <= 20 for 2 consecutive periods of 60 seconds for the metric dimensions:

Remove 1 capacity units

30 seconds before allowing another scaling activity

☐

scale-out-high-CPU-asp

Simple scaling

Enabled

high-CPU-alarm

breaches the alarm threshold: CPUUtilization >= 40 for 2 consecutive periods of 60 seconds for the metric dimensions:

Add 1 capacity units

30 seconds before allowing another scaling activity

☐

Screenshot 5 - 1: Auto-scaled instances

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone
<input type="checkbox"/>	Autoscaled Instance 01:04:46	i-0e78b765f62f58db3	Running	t2.nano	2/2 checks passed	View alarms	us-east-1b
<input type="checkbox"/>	Autoscaled Instance 02:04:01	i-0fc27766dd631ab21	Running	t2.nano	Initializing	View alarms	us-east-1a

Step 6: Test Traffic

See `generate_test_traffic.sh`

Test traffic was generated to the load balancer via a script which prompts the user for the DNS name of the load balancer. The traffic is generated using `curl` to send 100 HTTP requests to the load balancer using a loop.

Screenshot 6: Test traffic generation script

```
# Prompt user to enter the DNS name of the load balancer
read -p "Enter the DNS name of the load balancer: " LB_DNS_NAME

# Loop to send requests with different URLs, supressing output
for i in {1..100}; do
    curl -s "http://$LB_DNS_NAME/$i" > /dev/null &
done
```

Step 7: Load Distribution

See load_logs.sh

This script is to be used after generating test traffic to the load balancer. It prompts the user for the ip address of the auto-scaled instance, then shows the Apache access log.

Due to lack of time, here is the browser load distribution.

Screenshot 7: Load Distribution (browser)



Step 8: Custom metrics

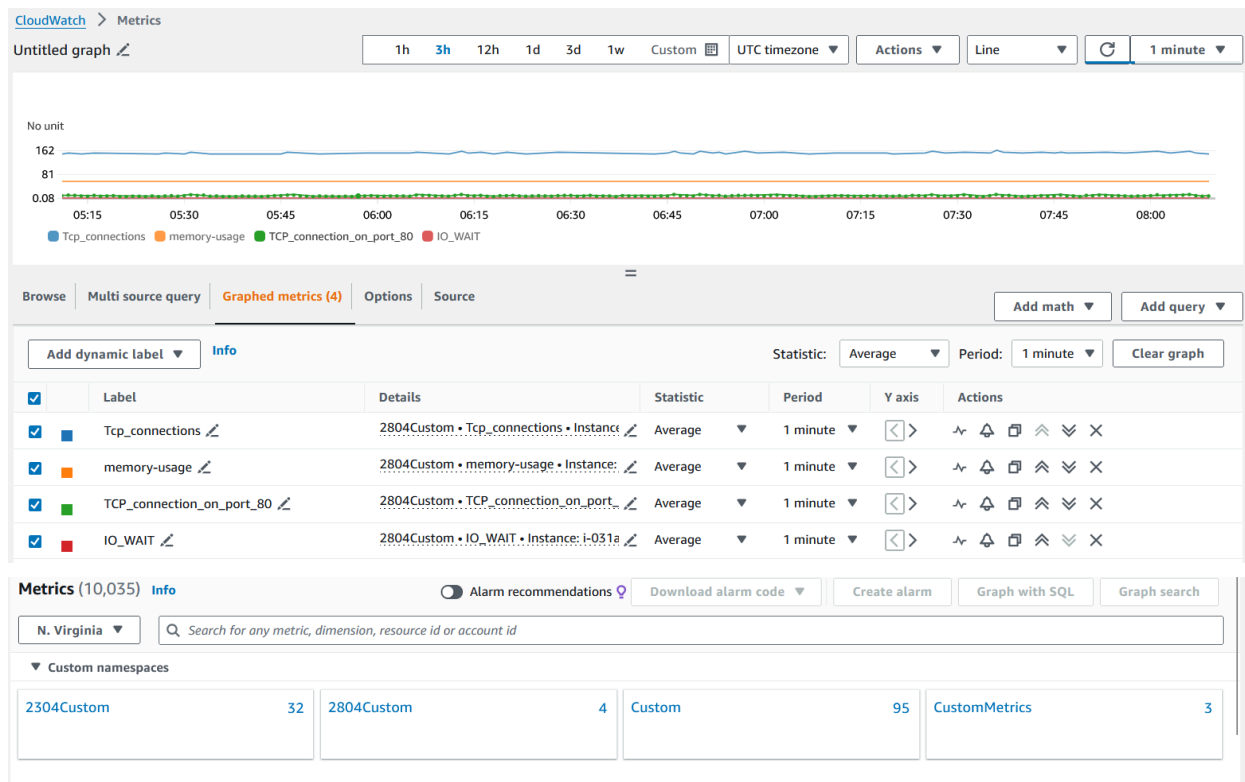
See `push_metrics.sh`

This script is added to the user data of the launch template for auto-scaled instances. The script collects performance and health metrics on those instances, and then pushes those metrics to CloudWatch every minute using a cron job. This script failed and I used SSH to add the custom metrics script to the AMI.

The metrics collected are:


1. **Used Memory:** This metric ensures an instance has sufficient memory to efficiently handle its workload.
2. **TCP Connections:** It measures the total number of TCP connections to the instance, reflecting network activity. High TCP connections indicate increased traffic.
3. **TCP Connections from Port 80:** Specifically tracks incoming HTTP traffic.
4. **IO Wait Time:** This represents the percentage of time the CPU spends waiting for I/O operations to complete.
5. **Process Count:** It counts the total number of processes running on the instance, reflecting system activity.

Screenshot 8: Custom metrics in CloudWatch



Additional Functionality

- Capture your configuration using your own customised Terraform script:
<https://github.com/MartyRath/aws-web-app-autoscaling>
- Screenshot: HTTPS set up:



- Web app communicates with Mongo database in private subnet.
- Failed attempt - Deploying app via user_data script. See deploy_app.sh